Checkers Is Solved

Jonathan Schaeffer,* Neil Burch, Yngvi Björnsson,† Akihiro Kishimoto,‡ Martin Müller, Robert Lake, Paul Lu, Steve Sutphen

The game of checkers has roughly 500 billion billion possible positions (5×10^{20}). The task of solving the game, determining the final result in a game with no mistakes made by either player, is daunting. Since 1989, almost continuously, dozens of computers have been working on solving checkers, applying state-of-the-art artificial intelligence techniques to the proving process. This paper announces that checkers is now solved: Perfect play by both sides leads to a draw. This is the most challenging popular game to be solved to date, roughly one million times as complex as Connect Four. Artificial intelligence technology has been used to generate strong heuristic-based game-playing programs, such as Deep Blue for chess. Solving a game takes this to the next level by replacing the heuristics with perfection.

Since Claude Shannon's seminal paper on the structure of a chess-playing program in 1950 (1), artificial intelligence researchers have developed programs capable of challenging and defeating the strongest human players in the world. Superhuman-strength programs exist for popular games such as chess [Deep Fritz (2)], checkers [Chinook (3)], Othello [Logistello (4)], and Scrabble [Maven (5)]. However strong these programs are, they are not perfect. Perfection implies solving a game—determining the final result (game-theoretic value) when neither player makes a mistake. There are three levels of solving a game (6). For the lowest level, ultraweakly solved, the perfect-play result, but not a strategy for achieving that value, is known [e.g., in Hex the first player wins, but for large board sizes the winning strategy is not known (7)]. For weakly solved games, both the result and a strategy for achieving it from the start of the game are known [e.g., in Go Moku the first player wins and a program can demonstrate the win (6)]. Strongly solved games have the result computed for all possible positions that can arise in the game [e.g., Awari (8)].

Checkers (8 × 8 draughts) is a popular game enjoyed by millions of people worldwide, with many annual tournaments and a series of competitions that determine the world champion. There are numerous variants of the game played around the world. The game that is popular in North America and the (former) British Commonwealth has pieces (checkers) moving forward one square diagonally, kings moving forward or backward one square diagonally, and a forced-capture rule [see supporting online material (SOM) text].

Department of Computing Science, University of Alberta, Edmonton, Alberta T6G 2E8, Canada.

*To whom correspondence should be addressed. E-mail: jonathan@cs.ualberta.ca
†Present address: Department of Computer Science, Reykjavik University, Kringlan 1, IS-103, Iceland.
‡Present address: Department of Media Architecture, Future University, Hakodate, 116-2 Kamedanakano-cho Hakodate Hokkaido, 041-8655, Japan.

The effort to solve checkers began in 1989, and the computations needed to achieve that result have been running almost continuously since then. At the peak in 1992, more than 200 processors were devoted to the problem simultaneously. The end result is one of the longest running computations completed to date.

With this paper, we announce that checkers has been weakly solved. From the starting position (Fig. 1, top), we have a computational proof that checkers is a draw. The proof consists of an explicit strategy that never loses—the program can achieve at least a draw against any opponent, playing either the black or white pieces. That checkers is a draw is not a surprise; grandmaster players have conjectured this for decades.

The checkers result pushes the boundary of artificial intelligence (AI). In the early days of AI research, the easiest path to achieving high performance was believed to be emulating the human approach. This was fraught with difficulty, especially the problems of capturing and encoding human knowledge. Human-like strategies are not necessarily the best computational strategies. Perhaps the biggest contribution of applying AI technology to developing game-playing programs was the realization that a search-intensive ("brute-force") approach could produce high-quality performance using minimal application-dependent knowledge. Over the past two decades, powerful search techniques have been developed and successfully applied to problems such as optimization, planning, and bioinformatics. The checkers proof extends this approach by developing a program that has little need for application-dependent knowledge and is almost completely reliant on search. With advanced AI algorithms and improved hardware (faster processors, larger memories, and larger disks), it has become possible to push the limits on the type and size of problems that can be solved. Even so, the checkers search space (5×10^{20}) represents a daunting challenge for today's technology.

Computer proofs in areas other than games have been done numerous times. Perhaps the

best known is the four-color theorem (9). This deceptively simple conjecture—that given an arbitrary map with countries, you need at most four different colors to guarantee that no two adjoining countries have the same color—has been extremely difficult to prove analytically. In 1976, a computational proof was demonstrated. Despite the convincing result, some mathematicians were skeptical, distrusting proofs that had not been verified using human-derived theorems. Although important components of the checkers proof have been independently verified, there may be skeptics.

This article describes the background behind the effort to solve checkers, the methods used for achieving the result, an argument that the result is correct, and the implications of this research. The computer proof is online (10).

Background. The development of a strong checkers program began in the 1950s with Arthur Samuel's pioneering work in machine learning. In 1963, his program played a match against a capable player, winning a single game. This result was heralded as a triumph for the fledgling field of AI. Over time, the result was exaggerated, resulting in claims that checkers was now "solved" (3).

The Chinook project began in 1989 with the goal of building a program capable of challenging the world checkers champion. In 1990, Chinook earned the right to play for the World Championship. In 1992, World Champion Marion Tinsley narrowly defeated Chinook in the title match. In the 1994 rematch, Tinsley withdrew part way due to illness. He passed away eight months later. By 1996 Chinook was much

Fig. 1. Black to play and draw. (**Top**) Standard starting board. (**Bottom**) Square numbers used for move notation.

stronger than all human players, and with faster processors this gap has only grown (*3*).

Tinsley was the greatest checkers player that ever lived, compiling an incredible record that included only three losses in the period from 1950 to 1991. The unfinished Tinsley match left the question unanswered as to who was the better player. If checkers were a proven draw, then a "perfect" Chinook would never lose. As great as Tinsley was, he did occasionally make losing oversights. Hence, solving checkers would once and for all establish computers as better checkers players than all (fallible) humans.

Numerous nontrivial games have been solved, including Connect Four (*6, 11*), Qubic (*6*), Go-Moku (*6*), Nine Men's Morris (*12*), and Awari (*8*). The perfect-play result and a strategy for achieving that result is known for these games. How difficult is it to solve a game? There are two dimensions to consider (*6*): (i) decision complexity, the difficulty of making correct move decisions, and (ii) space complexity, the size of the search space.

Checkers is considered to have high decision complexity (it requires extensive skill to make strong move choices) and moderate space complexity (5×10^{20}) (Table 1). All the games solved thus far have either low decision complexity (Qubic and Go-Moku), low space complexity (Nine Men's Morris, size 10^{11},

and Awari, size 10^{12}), or both (Connect Four, size 10^{14}).

Solving checkers. Checkers represents the most computationally challenging game solved to date. The proof procedure has three algorithm/data components (*13*): (i) Endgame databases (backward search). Computations from the end of the game back toward the starting position have resulted in a database of 3.9×10^{13} positions (all positions with ≤10 pieces on the board) for which the game-theoretic value has been computed (strongly solved). (ii) Proof-tree manager (forward search). This component maintains a tree of the proof in progress (a sequence of moves and their best responses), traverses it, and generates positions that need to be explored to further the proof's progress. (iii) Proof solver (forward search). Given a position to search by the manager, this component uses two programs to determine the value of the position. These programs approach the task in different ways, thus increasing the chances of obtaining a useful result. Figure 2 shows the forward and backward search interactions in the checkers search space.

In the manager, the proof tree can be hand-seeded with an initial line of play. From the literature (*14*), a single "best" line of play was identified and used to guide the initial foray of the manager into the depths of the search tree. Although not essential for the proof, this is an

important performance enhancement. It allows the proof process to immediately focus its work on the parts of the search space that are likely to be relevant. Without it, the manager may spend unnecessary effort looking for an important line to explore. The line leads from the start of the game into the endgame databases (Fig. 2).

Backward search. Positions at the end of the game can be searched and their win/loss/draw value determined. The technique is called retrograde analysis and has been successfully used for many games. The algorithm works backward by starting at the end of the game and working toward the start. It enumerates all one-piece positions, determining their value (in this case, a trivial win for the side with the piece). Next, all two-piece positions are enumerated and analyzed. The analysis for each position eventually leads to a one-piece position with a known value, or a repeated position (draw). Next, all the three-piece positions are tackled, and so forth (SOM text). Our program has computed all the positions with ≤10 pieces on the board. The endgame databases are crucial to solving checkers. The checkers forced-capture rule quickly results in many pieces being removed from the board, giving rise to a position with ≤10 pieces—and a known value.

The databases contain the win/loss/draw result for a position, not the number of moves to a win/

Fig. 2. Forward and backward search. The number of pieces on the board are plotted (vertically) versus the logarithm of the number of positions (Table 1). The shaded area shows the endgame database part of the proof—i.e., all positions with ≤10 pieces. The inner oval area shows that only a portion of the search space is relevant to the proof. Positions may be irrelevant because they are unreachable or are not required for the proof. The small open circles indicate positions with more than 10 pieces for which a value has been proven by a solver. The dotted line shows the boundary between the top of the proof tree that the manager sees (and stores on disk) and the parts that are computed by the solvers (and are not saved in order to reduce disk storage needs). The solid seeded line shows a "best" sequence of moves.

Table 1. The number of positions in the game of checkers. For example, the possible positions for one piece include 32 squares for the Black king, 32 squares for the White king, 28 squares for a Black checker, and 28 squares for a White checker, for a total of 120 positions.

Pieces	Number of positions
1	120
2	6,972
3	261,224
4	7,092,774
5	148,688,232
6	2,503,611,964
7	34,779,531,480
8	406,309,208,481
9	4,048,627,642,976
10	34,778,882,769,216
Total 1–10	39,271,258,813,439
11	259,669,578,902,016
12	1,695,618,078,654,976
13	9,726,900,031,328,256
14	49,134,911,067,979,776
15	218,511,510,918,189,056
16	852,888,183,557,922,816
17	2,905,162,728,973,680,640
18	8,568,043,414,939,516,928
19	21,661,954,506,100,113,408
20	46,352,957,062,510,379,008
21	82,459,728,874,435,248,128
22	118,435,747,136,817,856,512
23	129,406,908,049,181,900,800
24	90,072,726,844,888,186,880
Total 1–24	500,995,484,682,338,672,639

loss. Independent research has discovered a 10-piece database position requiring a 279-ply move sequence to demonstrate a forced win (a ply is one move by one player) (15). This is a conservative bound; the win length has not been computed for the more difficult (and more interesting) database positions.

The complete 10-piece databases contain 39 trillion positions (Table 1). They are compressed into 237 gigabytes, an average of 154 positions per byte. A custom compression algorithm was used that allows for rapid localized real-time decompression (16). This means that the backward and forward search programs can quickly extract information from the databases.

The first databases, constructed in 1989, were for less than or equal to four pieces. In 1994, Chinook used a subset of the eight-piece database for the Tinsley match (3). By 1996, the eight-piece database was completed, giving rise to hope that checkers could be solved. However, the problem was still too hard, and the effort came to a halt. In 2001, computer capabilities had increased substantially, and the effort was restarted. It took 7 years (1989 to 1996) to compute the original eight-piece databases; in 2001 it took only a month. In 2005, the 10-piece database computation finished. At this point, all computational resources were focused on the forward search effort.

Forward search. Development of the forward search program began in 2001, with the production version up and running in 2004. The forward search consists of two parts: the proof-tree manager, which builds the proof by identify-

ing positions that need to be assessed, and the proof solvers, which search individual positions.

The manager maintains the master copy of the proof and uses the Proof Number search algorithm (6) to identify a prioritized list of positions that need to be examined. Typically, several hundred positions of interest are generated at a time so as to keep multiple computers busy. Over the past year, we used an average of 50 computers simultaneously.

The solvers get a position to evaluate from the manager. The result of a position evaluation can be proven (win, loss, or draw), partially proven (at least a draw, at most a draw), or heuristic (an estimate of how good or bad a position is). Proven positions need no further work; partially proven positions need additional work if the manager determines that a proven value is needed. If no proven information is available then the solver returns a heuristic assessment of the position. The manager uses this assessment to prioritize which positions to consider next. The manager updates the proof tree with the new information, decides which positions need further investigation, and generates new work to do. This process is repeated until a proven result for the game is determined.

The solver uses two search programs to evaluate a position. The first program (targeted at 15 s, but sometimes much longer) uses Chinook to determine a heuristic value for the position (alpha-beta search to nominal search depths of 17 to 23 ply). Occasionally, this search determines that the position is a proven win or loss. Chinook was not designed to produce a

proven draw, only a heuristic draw; demonstrating proven draws in a heuristic search seriously degrades performance.

The alpha-beta search algorithm is the mainstay of game-playing programs. The algorithm does a depth-first, left-to-right traversal of the search tree (17) (SOM text). The algorithm propagates heuristic bounds on the value of a position: the minimum value that the side about to move can achieve and the maximum value that the side about to move can be limited to by the opponent. Lines of play that are provably outside this range are irrelevant and can be eliminated (cut off). A d-ply search with an average of b moves to consider in every position results in a tree with roughly b^d positions. In the best case, the alpha-beta algorithm only needs to examine roughly $b^{d/2}$ positions (16).

If Chinook does not find a proven result, then a second program is invoked (100 s). It uses the Df-pn algorithm (18), a space-efficient variant of Proof Number search. The search returns a proven, partially proven, or unknown result.

Algorithms based on proof numbers maintain a measure of the difficulty of proving a position. This difficulty is expressed as a proof number, a lower bound on the minimum number of positions that need to be explored to result in the position being proven. The algorithm repeatedly expands the tree below the position requiring the least effort to affect the original position (a "best-first" approach). The result of that search is propagated back up the tree, and a new best candidate to consider is determined. Proof number search was specifically invented to facilitate the proving of games. The Df-pn variant builds the search tree in a depth-first manner, requiring less computer storage.

Iterative search algorithms are commonplace in the AI literature. Most iterate on search depth (first 1 ply, then 2, then 3, etc.). The manager uses the new approach of iterating on the error in Chinook's heuristic scores (13). The manager uses a threshold, t, and temporarily assumes that all heuristic scores $\geq t$ are wins and all scores $\leq -t$ are losses. It then proves the result given this assumption. Once completed, t is increased to $t+\Delta$, and the process is repeated. Eventually t reaches the value of a win and the proof is complete. This iterative approach concentrates the effort on forming the outline of the proof with low values of t, and then fleshing out the details with the rest of the computation.

One complication is the graph-history interaction (GHI) problem. It is possible to reach the same position through two different sequences of moves. This means that some draws depend on the moves played leading to the duplicated position. In standard search algorithms, GHI may cause some positions to be incorrectly inferred as draws. Part of this research project was to develop an improved algorithm for addressing the GHI problem (19).

Correctness. Given a computation that has run for so long on many processors, an important

Table 2. Openings solved. Shown are the opening moves (using the standard square number scheme in Fig. 1, bottom), the result, the number of positions given to the solvers, and the position farthest from the start of the game that was searched (Max ply). The last two columns give the size and ply depth of the pruned minimal proof tree. Note that the total does not match the sum of the 19 openings. The combined tree has some duplicated nodes, which have been removed when reporting the total.

No.	Opening	Proof	Searches	Max ply	Minimal size	Max ply
1	09-13 22-17 13-22	Draw	736,984	56	275,097	55
2	09-13 21-17 05-09	Draw	1,987,856	154	684,403	85
3	09-13 22-18 10-15	Draw	715,280	103	265,745	58
4	09-13 23-18 05-09	Draw	671,948	119	274,376	94
5	09-13-23-19 11-16	Draw	964,193	85	358,544	71
6	09-13 24-19 11-15	Draw	554,265	53	212,217	49
7	09-13 24-20 11-15	Draw	1,058,328	59	339,562	58
8	09-14 23-18 14-23	≤Draw	2,202,533	77	573,735	75
9	10-14 23-18 14-23	≤Draw	1,296,790	58	336,175	55
10	10-15 22-18 15-22	≤Draw	543,603	60	104,882	41
11	11-15 22-18 15-22	≤Draw	919,594	67	301,310	59
12	11-16 23-19 16-23	≤Draw	1,969,641	69	565,202	64
13	12-16 24-19 09-13	Loss	205,385	44	49,593	40
14	12-16 24-19 09-14	≤Draw	61,279	45	23,396	44
15	12-16 24-19 10-14	≤Draw	21,328	31	8,917	31
16	12-16 24-19 10-15	≤Draw	31,473	35	13,465	35
17	12-16 24-19 11-15	≤Draw	23,803	34	9,730	34
18	12-16 24-19 16-20	≤Draw	283,353	49	113,210	49
19	12-16 24-19 08-12	≤Draw	266,924	49	107,109	49
Overall		Draw	Total	Max	Total	Max
			15,123,711	154	3,301,807	94

question to ask is "Are the results correct?" Early on in the computation, we realized that there were many potential sources of errors, including algorithm bugs and data transmission errors. Great care has been taken to eliminate any possibility of error by verifying all computation results and doing consistency checks. As well, some of the computations have been independently verified (SOM text).

Even if an error has crept into the calculations, it likely does not change the final result. Assume a position that is 40 ply away from the start is incorrect. The probability that this erroneous result can propagate up 40 ply and change the value for the game of checkers is vanishingly small (20).

Results. Our approach to solving the game was to determine the game-theoretic result by doing the least amount of work. In tournament checkers, the standard starting position (Fig. 1, top) is considered "boring," so the first three moves (ply) of a game are randomly chosen at the start. The checkers proof consisted of solving 19 three-move openings, leading to a determination of the starting position's value: a draw. Although there are roughly 300 three-move openings, more than 100 are duplicates (move transpositions). The rest can be proven to be irrelevant by an alpha-beta search.

Table 2 shows the results for the 19 openings solved to determine the perfect-play result for checkers. (Other openings have been solved but are not included here.) After an opening was proven, a postprocessing program pruned the tree to eliminate all the computations that were not part of the smallest proof tree. In hindsight, the pruned work was unnecessary, but it was not so at the time when it was assigned for evaluation. Figure 3 shows the proof tree for the first 3 ply.

The leftmost move sequence in Fig. 3 is as follows: Black moves from 09 to 13 (represented using the standard checkers notation 09-13), White replies with 22-17, and then Black moves 13-22. The resulting position has been searched and shown to be a draw (opening line 1 in Fig. 3).

That means the position after 22-17 is also a draw, given that there is only one legal move available (13-22) and it is a proven draw. What is the value of the position after Black moves 09-13? To determine this, all possible moves for White have to be considered. The move 22-17 guarantees White at least a draw (at most a draw for Black). But it is possible that this position is a win for White (and a loss for Black). The remaining moves (21-17, 22-18, 23-18, 23-19, 24-19, and 24-20; opening lines 2 to 7 in Fig. 3) are all shown to be at least a draw for Black. Hence, White prefers the move 22-17 (no worse than any other move). Thus, 09-13 leads to a draw (White will move 22-17 in response).

Given that 09-13 is a draw, it remains to demonstrate that the other opening moves cannot win for Black. Note that some openings have a proven result, whereas for others only the partial result that was necessary for the proof was computed. The number of openings is small because the forced-capture rule was exploited. Opening lines 13 to 19 in Fig. 3 are needed to prove that the opening 12-16 is not a win. Actually, one opening would have sufficed (12-16, 23-19, and 16-23). However, human analysts consider this line to be a win for Black, and the preliminary analysis agreed. Hence, the seven openings beginning with the moves 12-16 and 24-19 were proven instead. This led to the least amount of computing.

There is anecdotal evidence that the proof tree is correct. Main lines of play were manually compared to human analysis (14), with no errors found in the computer's results (unimportant errors were found in the human analysis).

The proof tree shows the perfect lines of play needed to achieve a draw. If one side makes a losing mistake, the proof tree may not necessarily show how to win. This additional information is not necessary for proving the draw result.

The stored proof tree is only 10^7 positions. Saving the entire proof tree, from the start of the game so that every line ends in an endgame database position, would require many tens of

terabytes, resources that were not available. Instead, only the top of the proof tree, the information maintained by the manager, is stored on disk. When a user queries the proof, if the end of a line of play in the proof is reached, then the solver is used to continue the line into the databases. This substantially reduces the storage needs, at the cost of recomputing (roughly 2 min per search).

The longest line analyzed was 154 ply. The position at the end of this line was analyzed by the solver, and that analysis may have gone 20 or more ply deep. At the end of this analysis is a database position, which could be the result of several hundred ply of analysis. This provides supporting evidence of the difficulty of checkers—for computers and humans.

How much computation was done in the proof? Roughly speaking, there are 10^7 positions in the stored proof tree, each representing a search of 10^7 positions (relatively small because of the extensive disk operations). Hence, 10^{14} is a good ballpark estimate of the forward search effort.

Should we be impressed with "only" 10^{14} computations? At one extreme, checkers could be solved using storage—build endgame databases for the complete search space. This would require 5×10^{20} data entries. Even an excellent compression algorithm might only reduce this to 10^{18} bytes, impractical with today's technology. This also makes it unlikely that checkers will soon be strongly solved.

An alternative would be to use only computing—i.e., build a search tree using the alpha-beta algorithm. Consider the following unreasonably optimistic assumptions: number of moves to consider is eight in noncapture positions, a game lasts 70 ply, all captures are of a single piece (23 capture moves), and the alpha-beta search does the least possible work. The assumptions result in a search tree of $8^{(70-23)} = 8^{47}$ states. The perfect alpha-beta search will halve the exponent, leading to a search of roughly $8^{47/2} \approx 10^{24}$. This would take more than a lifetime to search, given current technology.

Conclusion. What is the scientific significance of this result? The early research was devoted to developing Chinook and demonstrating superhuman play in checkers, a milestone that predated the Deep Blue success in chess. The project has been a marriage of research in AI and parallel computing, with contributions made in both of these areas. This research has been used by a bioinformatics company; real-time access of very large data sets for use in parallel search is as relevant for solving a game as it is for biological computations.

The checkers computation pushes the boundary of what can be achieved by search-intensive algorithms. It provides compelling evidence of the power of limited-knowledge approaches to artificial intelligence. Deep search implicitly uncovers knowledge. Furthermore, search algorithms are well poised to take advantage of the

Fig. 3. The first three moves of the checkers proof tree. Move sequences are indicated using the notation from Fig. 1B, with the from-square and to-square of the move separated by a hyphen. The result of each position is given for Black, the first player to move (=D, a proven draw; =L, a proven loss; <=D, loss or draw; and >=D, draw or win). In some positions, only one move needs to be considered; the rest are cut off, as indicated by the rotated "T". Some positions have only one legal move because of the forced-capture rule.

increase in on-chip parallelism that multicore computing will soon offer. Search-intensive approaches to AI will play an increasingly important role in the evolution of the field.

With checkers finished, the obvious question is whether chess is solvable. Checkers has roughly the square root of the number of positions in chess (somewhere in the 10^{40} to 10^{50} range). Given the effort required to solve checkers, chess will remain unsolved for a long time, barring the invention of new technology. The disk-flipping game of Othello is the next popular game that is likely to be solved, but it will require considerably more resources than were needed to solve checkers (7).

References and Notes
1. C. Shannon, *Philos. Mag.* **41**, 256 (1950).
2. "The Duel: Man vs. Machine" (www.rag.de/microsite_chess_com).
3. J. Schaeffer, *One Jump Ahead* (Springer-Verlag, New York, 1997).
4. M. Buro, *IEEE Intell. Syst. J.* **14**, 12 (1999).
5. B. Sheppard, thesis, Universiteit Maastricht, Maastricht, Netherlands (2002).
6. V. Allis, thesis, University of Limburg, Maastricht, Netherlands (1994).
7. J. van den Herik, J. Uiterwijk, J. van Rijswijck, *Artif. Intell.* **134**, 277 (2002).
8. J. Romein, H. Bal, *IEEE Computer* **36**, 26 (2003).
9. K. Appel, W. Haken, *Sci. Am.* **237**, 108 (1977).
10. Chinook Web site, with proof (www.cs.ualberta.ca/~chinook).
11. John's Connect Four Playground (http://homepages.cwi.nl/~tromp/c4/c4.html).
12. R. Gasser, thesis, ETH Zürich, Switzerland (1995).
13. J. Schaeffer *et al.*, "Solving Checkers" (www.ijcai.org/papers/0515.pdf).
14. R. Fortman, *Basic Checkers* (http://home.clara.net/davey/basicche.html).
15. "Longest 10PC MTC" (http://pages.prodigy.net/eyg/Checkers/longest-10pc-mtc.htm).
16. J. Schaeffer *et al.*, in *Advances in Computer Games*, J. van den Herik, H. Iida, E. Heinz, Eds. (Kluwer, Dordrecht, Netherlands, 2003), pp. 193–210.
17. D. Knuth, R. Moore, *Artif. Intell.* **6**, 293 (1975).
18. A. Nagai, thesis, University of Tokyo, Japan (2002).
19. A. Kishimoto, M. Müller, in *Proceedings of the Nineteenth National Conference on Artificial Intelligence* (AAAI Press, Menlo Park, CA, 2004) pp. 644–649.
20. D. Beal, thesis, Universiteit Maastricht, Maastricht, Netherlands (1999).
21. The support of Canada's Natural Sciences and Engineering Research Council (NSERC), Alberta's Informatics Circle of Research Excellence (iCORE), and the Canada Foundation for Innovation is greatly appreciated. Numerous people contributed to this work, including M. Bryant, J. Culberson, B. Gorda, B. Knight, D. Szafron, K. Thompson, and N. Treloar.

TLR3 Deficiency in Patients with Herpes Simplex Encephalitis

Shen-Ying Zhang,[1,2,3] Emmanuelle Jouanguy,[1,2,3] Sophie Ugolini,[4] Asma Smahi,[5] Gaëlle Elain,[6] Pedro Romero,[7] David Segal,[8] Vanessa Sancho-Shimizu,[1,2] Lazaro Lorenzo,[1,2] Anne Puel,[1,2] Capucine Picard,[1,2,9] Ariane Chapgier,[1,2] Sabine Plancoulaine,[1,2] Matthias Titeux,[10] Céline Cognet,[4] Horst von Bernuth,[1,2] Cheng-Lung Ku,[1,2] Armanda Casrouge,[1,2] Xin-Xin Zhang,[3] Luis Barreiro,[11] Joshua Leonard,[8] Claire Hamilton,[1,2] Pierre Lebon,[12] Bénédicte Héron,[13] Louis Vallée,[14] Lluis Quintana-Murci,[11] Alain Hovnanian,[10] Flore Rozenberg,[12] Eric Vivier,[4] Frédéric Geissmann,[6] Marc Tardieu,[15] Laurent Abel,[1,2] Jean-Laurent Casanova[1,2,3,16]*

Some Toll and Toll-like receptors (TLRs) provide immunity to experimental infections in animal models, but their contribution to host defense in natural ecosystems is unknown. We report a dominant-negative *TLR3* allele in otherwise healthy children with herpes simplex virus 1 (HSV-1) encephalitis. TLR3 is expressed in the central nervous system (CNS), where it is required to control HSV-1, which spreads from the epithelium to the CNS via cranial nerves. TLR3 is also expressed in epithelial and dendritic cells, which apparently use TLR3-independent pathways to prevent further dissemination of HSV-1 and to provide resistance to other pathogens in TLR3-deficient patients. Human TLR3 appears to be redundant in host defense to most microbes but is vital for natural immunity to HSV-1 in the CNS, which suggests that neurotropic viruses have contributed to the evolutionary maintenance of TLR3.

The contribution of Toll and Toll-like receptors to immunity has been studied extensively in the past decade. Toll-deficient *Drosophila* were shown to be susceptible to experimental infections with certain fungi in 1996 (1), and a Toll-like receptor 4 (TLR4) null mutation in mice resistant to lipopolysaccharide (LPS) but susceptible to certain Gram-negative bacteria was identified in 1998 (2). Mice deficient for individual TLRs have since been generated and shown to have diverse infectious phenotypes, from susceptibility to resistance, depending on the TLR-pathogen combination (3). However, it remains unclear whether TLRs play nonredundant roles—beneficial or detrimental—in natural, as opposed to experimental, infections. This biological question is important, because

natural selection acts on a given species in the setting of natural (rather than experimental) ecosystems. The human model is particularly suitable for analyses of the relevance of genes such as those of TLRs to host defense in natural ecosystems (4). Nevertheless, although many studies have suggested that TLR genes are involved in human infectious diseases, this has not been unambiguously demonstrated (5). In particular, no primary immunodeficiency involving TLRs has been identified.

The discovery of inherited interleukin 1 receptor-associated kinase-4 (IRAK-4) deficiency in children with bacterial diseases implicated human TLRs, interleukin-1 receptors (IL-1Rs), or both in host defense (6, 7). However, the narrow range of infections documented in such patients

indicates that IRAK-4–dependent, TLR-mediated immunity is redundant for protective immunity to most microbes. In particular, IRAK-4–deficient patients are not susceptible to herpes simplex virus 1 (HSV-1) encephalitis (HSE). In HSE, HSV-1 infects epithelial cells in the oral and nasal mucosa and progresses to the central nervous system (CNS) via the trigeminal or olfactory nerves (8). A genetic etiology of HSE was found in two children who lacked functional UNC-93B (9), an endoplasmic reticulum protein required for TLR3, TLR7, TLR8, and TLR9 signaling (10). Both UNC-93B– and IRAK-4–deficient patients fail to signal through TLR7, TLR8, and TLR9, but unlike IRAK-4–deficient patients (7), UNC-93B–deficient patients display impaired TLR3-dependent interferon-α (IFN-α) -β, and -λ production (9). Moreover, HSV-1 is a double-stranded DNA virus with double-stranded RNA (dsRNA) intermediates (11), and TLR3 recog-

[1]Human Genetics of Infectious Diseases, Institut National de la Santé et de la Recherche Médicale (INSERM), U550, Faculty Necker, Paris 75015, France. [2]University Paris René Descartes, Paris 75015, France. [3]French-Chinese Laboratory of Genetics and Life Sciences, Rui Jin Hospital, Shanghai Jiao Tong University, Shanghai 200025, China. [4]Marseille-Luminy Immunology Institute, Marseille 13288, France. [5]Department of Genetics, INSERM, U781, Necker Hospital, Paris 75015, France. [6]Laboratory of Mononuclear Cell Biology, INSERM, U838, Necker Hospital, Paris 75015, France. [7]Ludwig Institute for Cancer Research, Lausanne Branch, University Hospital, Lausanne 1005, Switzerland. [8]Experimental Immunology Branch, National Cancer Institute, National Institutes of Health, Bethesda, MD 20892, USA. [9]Center for the Study of Immunodeficiencies, Necker Hospital, Paris 75015, France. [10]INSERM, U563, University Toulouse Paul Sabatier, Toulouse 31000, France. [11]Centre National de la Recherche Scientifique, URA3012, Pasteur Institute, Paris 75015, France. [12]Virology, Cochin-Saint-Vincent de Paul Hospital, University Paris René Descartes, Paris 75014, France. [13]Pediatric Neurology, Trousseau Hospital, Paris 75012, France. [14]Pediatric Neurology, University Hospital, Lille 59037, France. [15]Pediatric Neurology, Bicêtre Hospital, University Paris Sud, Kremlin-Bicêtre 94270, France. [16]Pediatric Hematology-Immunology, Necker Hospital, Paris 75015, France.

*To whom correspondence should be addressed. E-mail: casanova@necker.fr

Science

Checkers Is Solved

Jonathan Schaeffer, Neil Burch, Yngvi Björnsson, Akihiro Kishimoto, Martin Müller, Robert Lake, Paul Lu, and Steve Sutphen

View the article online
https://www.science.org/doi/10.1126/science.1144079
Permissions
https://www.science.org/help/reprints-and-permissions

Use of this article is subject to the Terms of service